

```
model.compile(
    optimizer=keras.optimizers.RMSprop(learning_rate=0.0005),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

RMSprop
batch size=32_

להלן המטריקות החשובות מתוך הריצה מ-Epoch 1 עד Epoch 30:

• Epoch סופי (30):

• Training Accuracy = 0.8553

• Training Loss = 0.7190

• Validation Accuracy = 0.7956

• Validation Loss = 0.9029

• Learning Rate = 1.25×10^{-4}

• Best Validation לאורך כל האימון:

• (Epoch 30-ב) best Val Accuracy = 0.7956

• (Epoch 30-ב) best Val Loss = 0.9029

• Test Evaluation (בסוף):

• Test Accuracy = 0.7940

• Test Loss = 0.9163

בקיצור, המודל הגיע לכ-79.6% וולידציה עם לוס ≈ 0.90 , ו-79.4% טסט עם לוס ≈ 0.92 .

```
[ ] he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.Dropout(0.3),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),



    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        68,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Dense(num_classes, activation='softmax')
])
```



```
model.compile(
    optimizer=keras.optimizers.SGD(learning_rate=0.003, momentum=0.9),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

SGD+momentum

batch size=32_

להלן המטריקות החשובות מתוך הריצה מ־Epoch 1 עד Epoch 30:

• Epoch סופי (30):

Training Accuracy = 0.8531

Training Loss = 0.7654

Validation Accuracy = 0.7823

Validation Loss = 0.9931

Learning Rate = 3.75×10^{-4}

• Best Validation לאורך כל האימון:

Best Val Accuracy = 0.7867 (Epoch 28)



Best Val Loss = 0.9834 (Epoch 28)

• Test Evaluation (בסוף):

Test Accuracy = 0.7817

Test Loss = 1.0057

בקיצור, ה־Best Validation Accuracy הגיע לכ־78.7% עם $Loss \approx 0.98$, ואילו על ה־Test קיבלנו כ־78.17% עם $Loss \approx 1.01$.



```
[ ] he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.6),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.005)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        68,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.005)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.6),

    layers.Dense(num_classes, activation='softmax')
])
```

```
model.compile(
    optimizer=keras.optimizers.SGD(learning_rate=0.003, momentum=0.9, nesterov=True
),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

NAG

batch size=32

להלן המטריקות שחושבו מתוך הריצה (NAG, LR Schedule-) מ-Epoch 1 עד Epoch 30:

• Epoch סופי (30)

Training Accuracy = **0.8460**

Training Loss = **0.7593**

Validation Accuracy = **0.7789**

Validation Loss = **0.9922**

Learning Rate = 3.75×10^{-4}

• Best Validation לאורך כל האימון

Best Val Accuracy = **0.7867** (Epoch 27)

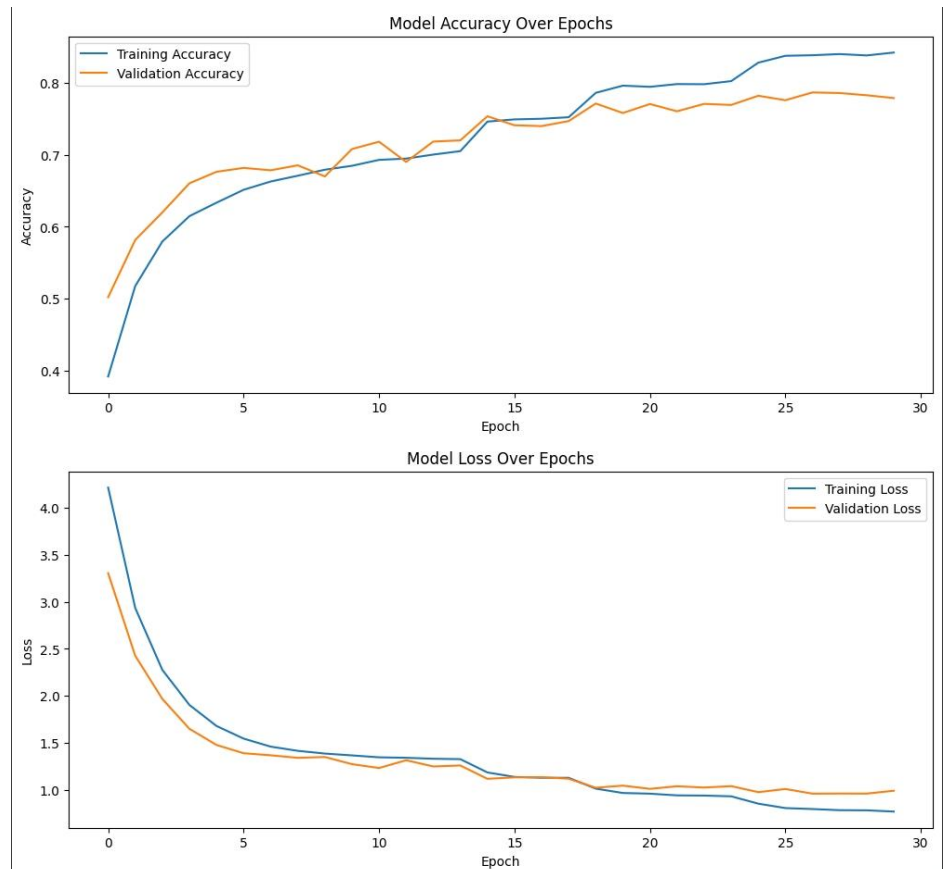
Best Val Loss = **0.9614** (Epoch 27)

• Test Evaluation (בסוף)

Test Accuracy = **0.7812**

Test Loss = **0.9785**

בקיצור – המודל הגיע ל-78.7% Best Validation Accuracy עם ≈ 0.96 Loss, ואילו על ה-Test קיבלנו כ-78.1% עם ≈ 0.98 Loss, פער קטן שמרמז על התאמה טובה וללא אובר-פיטינג משמעותי.



```
[18] he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        60, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Conv2D(
        60, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.6),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.005)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        60,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.005)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.6),

    layers.Dense(num_classes, activation='softmax')
])
```

```
model.compile(
    optimizer=keras.optimizers.Adagrad(learning_rate=0.01),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

AdaGrad

batch size=32

להלן המטריקות שחושבו מתוך הריצה (Adagrad, LR=0.01) מ־Epoch 1 עד Epoch 30:

• Epoch סופי (30):

Training Accuracy = 0.8219

Training Loss = 1.0225

Validation Accuracy = 0.7591

Validation Loss = 1.1991

Learning Rate = 1×10^{-2}

• Best Validation לאורך כל האימון:

Best Val Accuracy = 0.7625 (Epoch 28)

Best Val Loss = 1.1982 (Epoch 27)

• Test Evaluation (בסוף):

Test Accuracy = 0.7616

Test Loss = 1.2092

בקיצור – המודל הגיע ל־76.25% Val עם $\text{Loss} \approx 1.20$, ואילו על ה־Test קיבלנו כ־76.16% עם $\text{Loss} \approx 1.21$; התאמה סבירה וללא over-fitting משמעותי.

```
[18] he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.003)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.003)
    ),
    layers.BatchNormalization(),

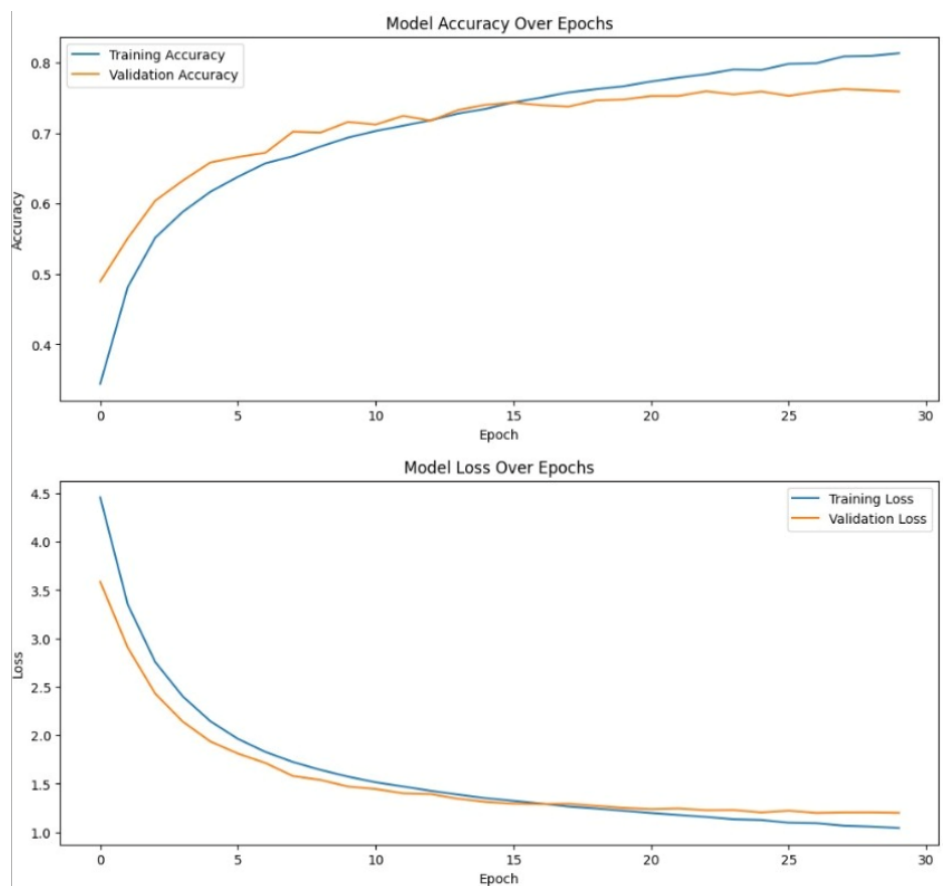
    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.6),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.005)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        68,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.L2(0.005)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.6),

    layers.Dense(num_classes, activation='softmax')
])
```




```
model.compile(
    optimizer=keras.optimizers.Adadelta(
        learning_rate=0.5
    ),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

AdaDelta

batch size=32

להלן המטריקות שחושבו מתוך הריצה (Adadelta, LR=0.5) מ-Epoch 1 עד Epoch 30:

• Epoch סופי (30):

Training Accuracy = 0.8095

Training Loss = 0.8889

Validation Accuracy = 0.7859

Validation Loss = 0.9328

Learning Rate = 1.25×10^{-1}

• Best Validation לאורך כל האימון:

Best Val Accuracy = 0.7859 (Epoch 30)

Best Val Loss = 0.9328 (Epoch 30)

• Test Evaluation (בסוף):

Test Accuracy = 0.7772

Test Loss = 0.9466

בקיצור – המודל הגיע ל-78.59% ב-Val עם ≈ 0.93 Loss, ואילו על ה-Test קיבלנו כ-77.72% עם ≈ 0.95 Loss; התאמה טובה וללא over-fitting משמעותי.

```
[ ] he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.004)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.004)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.004)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.35),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.004)
    ),
    layers.BatchNormalization(),

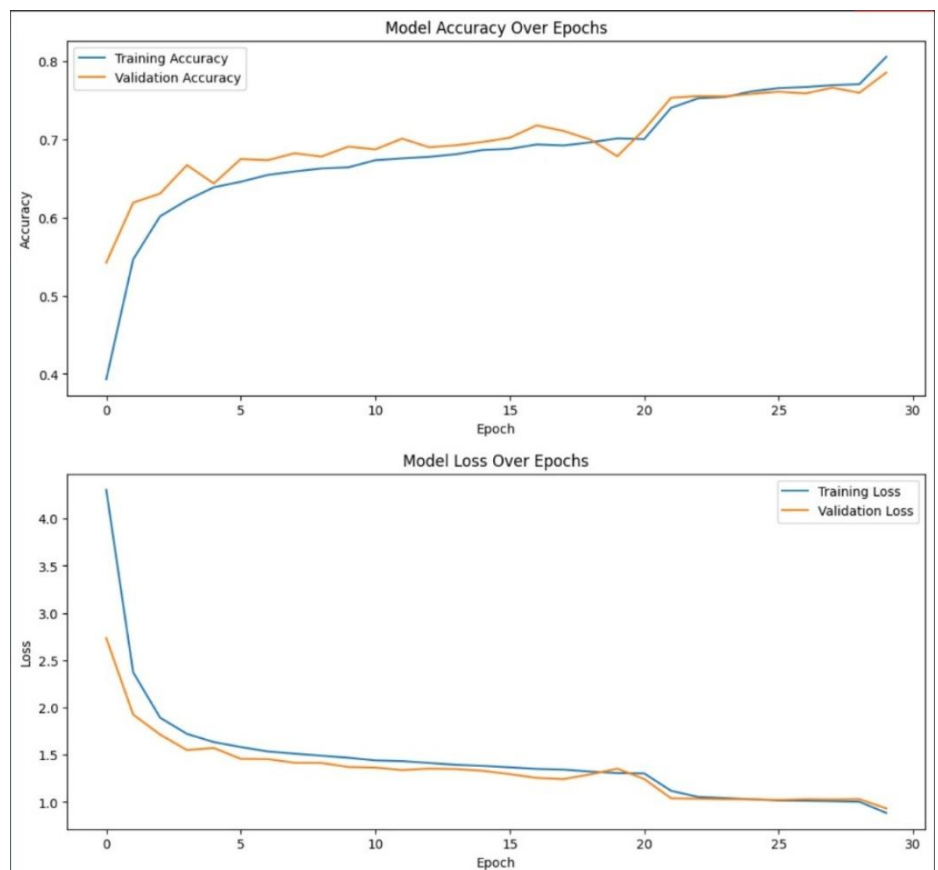
    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.004)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.006)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        68,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.006)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.45),

    layers.Dense(num_classes, activation='softmax')
])
```



```
model.compile(
    optimizer=tf.keras.optimizers.Adam(
        learning_rate=5e-4,
        beta_1=0.9,
        beta_2=0.999,
        epsilon=1e-7
    ),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

Adam

batch size=32

להלן המטריקות שחושבו מתוך הריצה מ-Epoch 1 עד Epoch 30 (LR Adam, התחלתי 5×10^{-4}):

• סופי Epoch (30):

Training Accuracy = 0.8389 –

Training Loss = 0.7868 –

Validation Accuracy = 0.7948 –

Validation Loss = 0.9226 –

Learning Rate = 1.25×10^{-4} –

• Best Validation לאורך כל האימון:

Best Val Accuracy = 0.7948 (Epoch 30) –

Best Val Loss = 0.9226 (Epoch 30) –

• Test Evaluation (בסוף):

Test Accuracy = 0.7902 –

Test Loss = 0.9310 –

בקיצור – המודל הגיע ל- $\approx 79.5\%$ Validation Accuracy ו- $\approx 79.0\%$ Test Accuracy, עם Validation Loss יורד ל-0.92;

התאמה טובה וללא over-fitting משמעותי.

```
he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Conv2D(
        68, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.35),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),



    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.002)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        68,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.003)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Dense(num_classes, activation='softmax')
])
```



```
model.compile(
    optimizer=tf.keras.optimizers.SGD(
        learning_rate=0.01
    ),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

SGD

batch size=32

להלן המטריקות החשובות שחושבו מתוך הריצה (SGD, LR=0.01) מ-Epoch 1 עד Epoch 30:

• Epoch סופי (30):

Training Accuracy = 0.8118 –

Training Loss = 0.8601 –

Validation Accuracy = 0.7516 –

Validation Loss = 1.0726 –

Learning Rate = 1×10^{-2} –

• Best Validation לאורך כל האימון:

Best Val Accuracy = 0.7516 (Epoch 30) –

Best Val Loss = 1.0726 (Epoch 30) –

• Test Evaluation (בסוף):

Test Accuracy = 0.7429 –

Test Loss = 1.1062 –

בקיצור, המודל הגיע לכ-75.16% דיוק בוולידציה ול-74.29% בדאטה של ה-Test, עם פער של כ-6% בין דיוק האימון לדיוק הוולידציה. נכון שהאימון יורד ללוס נמוך יחסית (0.86), אבל הוולידציה נותרת סביב 1.07, מה שמעיד על over-fitting מתון. ההתאמה בין וולידציה ל-Test טובה (פער $\approx 1\%$), ואין פה שיפור פתאומי שמעיד על רעש.

```
he_initializer = initializers.HeNormal()

model = keras.Sequential([
    layers.Input(shape=(32, 32, 3)),
    layers.BatchNormalization(),

    layers.Conv2D(
        20, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),

    layers.Conv2D(
        60, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Conv2D(
        60, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Conv2D(
        82, (3, 3), padding='same',
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.001)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Flatten(),

    layers.Dense(
        21,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.0005)
    ),
    layers.BatchNormalization(),

    layers.Dense(
        68,
        activation='leaky_relu',
        kernel_initializer=he_initializer,
        kernel_regularizer=regularizers.l2(0.0005)
    ),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Dense(num_classes, activation='softmax')
])
```

